

03_Deep_Learning_Development_with_PyTorch

October 7, 2023

1 Chapter 3 - Deep Learning Development with PyTorch

```
[ ]: import torch
import torchvision
```

1.1 Data Loading

```
[ ]: from torchvision.datasets import CIFAR10


train_data = CIFAR10(root="./train/",
                      train=True,
                      download=True)
```

Files already downloaded and verified

```
[ ]: # Accessing the Data
import matplotlib.pyplot as plt
plt.imshow(train_data[0][0])

#Accessing the label
print(train_data[0][1])
```

6



```
[ ]: test_data = CIFAR10(root="./test/",
                        train=False,
                        download=True)

print(len(test_data))

print(test_data.data.shape)
```

```
Files already downloaded and verified
10000
(10000, 32, 32, 3)
```

1.2 Data Transforms

```
[ ]: from torchvision import transforms

train_transforms = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize(
```

```

        (0.4914, 0.4822, 0.4465),
        (0.2023, 0.1994, 0.2010))])

train_data_transforms = CIFAR10(root="./train/",
                                train=True,
                                download=True,
                                transform=train_transforms)

```

Files already downloaded and verified

```

[ ]: data,label = train_data[0]
print(type(data))
print(data)

data_transforms,label_transforms = train_data_transforms[0]
print(type(data_transforms))
print(data_transforms.size())
print(data_transforms)

fig = plt.figure(figsize=(10,5))
fig.add_subplot(1,2,1)
plt.imshow(data)
fig.add_subplot(1,2,2)
plt.imshow(data_transforms.permute(1, 2, 0))

```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```

<class 'PIL.Image.Image'>
<PIL.Image.Image image mode=RGB size=32x32 at 0x14F3FDD23B0>
<class 'torch.Tensor'>
torch.Size([3, 32, 32])
tensor([[[[-2.4291, -2.4291, -2.4291, ..., -2.4291, -2.4291, -2.4291],
          [-2.4291, -2.4291, -2.4291, ..., -0.1223, -0.5293, -1.1109],
          [-2.4291, -2.4291, -2.4291, ..., -0.1029, -0.7232, -1.4404],
          ...,
          [-2.4291, -2.4291, -2.4291, ..., 0.3236, 0.5950, 0.9051],
          [-2.4291, -2.4291, -2.4291, ..., 0.8858, 1.1184, 1.2735],
          [-2.4291, -2.4291, -2.4291, ..., 1.4673, 1.4091, 1.3316]],

        [[[-2.4183, -2.4183, -2.4183, ..., -2.4183, -2.4183, -2.4183],
          [-2.4183, -2.4183, -2.4183, ..., -0.6286, -0.9826, -1.3562],
          [-2.4183, -2.4183, -2.4183, ..., -0.8056, -1.4152, -1.8873],
          ...,
          [-2.4183, -2.4183, -2.4183, ..., -0.3926, -0.0582, 0.3154],
          [-2.4183, -2.4183, -2.4183, ..., 0.2368, 0.4531, 0.6694],

```


```

[-2.4183, -2.4183, -2.4183, ..., 0.8071, 0.6891, 0.5908]],

[[-2.2214, -2.2214, -2.2214, ..., -2.2214, -2.2214, -2.2214],
 [-2.2214, -2.2214, -2.2214, ..., -0.9922, -1.2069, -1.4020],
 [-2.2214, -2.2214, -2.2214, ..., -1.3825, -1.8117, -2.0653],
 ...,
 [-2.2214, -2.2214, -2.2214, ..., -1.6361, -1.7141, -1.6751],
 [-2.2214, -2.2214, -2.2214, ..., -1.5971, -1.5580, -1.6946],
 [-2.2214, -2.2214, -2.2214, ..., -1.5190, -1.5580, -1.7336]]])

```

```
[ ]: <matplotlib.image.AxesImage at 0x14f3fdd36a0>
```



```

[ ]: test_transforms = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize(
        (0.4914, 0.4822, 0.4465),
        (0.2023, 0.1994, 0.2010))])

test_data_transforms = torchvision.datasets.CIFAR10(
    root="./test/",
    train=False,
    transform=test_transforms)

print(test_data)

```

Dataset CIFAR10

Number of datapoints: 10000

Root location: ./test/

Split: Test

1.3 Data Batching

```
[ ]: trainloader = torch.utils.data.DataLoader(  
        train_data_transforms,  
        batch_size=16,  
        shuffle=True)
```

```
[ ]: data_batch, labels_batch = next(iter(trainloader))  
print(data_batch.size())  
  
print(labels_batch.size())
```

```
torch.Size([16, 3, 32, 32])  
torch.Size([16])
```

```
[ ]: testloader = torch.utils.data.DataLoader(  
        test_data_transforms,  
        batch_size=16,  
        shuffle=False)
```

1.4 Model Design

1.4.1 Using Existing & Pre-trained models

```
[ ]: from torchvision import models  
  
vgg16 = models.vgg16(pretrained=True)
```

```
c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-  
packages\torchvision\models\_utils.py:208: UserWarning: The parameter  
'pretrained' is deprecated since 0.13 and may be removed in the future, please  
use 'weights' instead.  
  warnings.warn(  
c:\Users\kaasa\AppData\Local\Programs\Python\Python310\lib\site-  
packages\torchvision\models\_utils.py:223: UserWarning: Arguments other than a  
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed  
in the future. The current behavior is equivalent to passing  
`weights=VGG16_Weights.IMAGENET1K_V1`. You can also use  
`weights=VGG16_Weights.DEFAULT` to get the most up-to-date weights.  
  warnings.warn(msg)
```

```
[ ]: print(vgg16.features)  
  
print(vgg16.avgpool)  
  
print(vgg16.classifier)
```

```
Sequential(  
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
  (1): ReLU(inplace=True)
```

```

(2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(3): ReLU(inplace=True)
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(6): ReLU(inplace=True)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): ReLU(inplace=True)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace=True)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace=True)
(16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(18): ReLU(inplace=True)
(19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace=True)
(23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(25): ReLU(inplace=True)
(26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(27): ReLU(inplace=True)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace=True)
(30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
AdaptiveAvgPool2d(output_size=(7, 7))
Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace=True)
  (5): Dropout(p=0.5, inplace=False)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)

```

```
[ ]: import torch.nn as nn
vgg16.classifier[-1] = nn.Linear(4096,10)

print(vgg16.classifier)

device = "cuda" if torch.cuda.is_available() else "cpu"

vgg_model = vgg16.to(device = device)
```

```
Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace=True)
  (5): Dropout(p=0.5, inplace=False)
  (6): Linear(in_features=4096, out_features=10, bias=True)
)
```

1.5 The PyTorch NN Module (torch.nn)

```
[ ]: import torch.nn as nn
import torch.nn.functional as F

class SimpleNet(nn.Module):

    def __init__(self):
        super(SimpleNet, self).__init__()
        self.fc1 = nn.Linear(2048, 256)
        self.fc2 = nn.Linear(256, 64)
        self.fc3 = nn.Linear(64,2)

    def forward(self, x):
        x = x.view(-1, 2048)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.softmax(self.fc3(x),dim=1)
        return x
```

```
[ ]: simplenet = SimpleNet()
print(simplenet)

input = torch.rand(2048)
output = simplenet(input)
```

```
SimpleNet(
  (fc1): Linear(in_features=2048, out_features=256, bias=True)
  (fc2): Linear(in_features=256, out_features=64, bias=True)
```

```

        (fc3): Linear(in_features=64, out_features=2, bias=True)
    )

```

1.6 Training

```

[ ]: from torch import nn
import torch.nn.functional as F

class LeNet5(nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5) # <1>
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, int(x.nelement() / x.shape[0]))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

device = "cuda" if torch.cuda.is_available() else "cpu"
LeNet_model = LeNet5().to(device=device)

```

1.6.1 Fundamental Training Loop

Code Annotations:

- <1> Our training loop
- <2> Need to move inputs and labels to GPU if avail.
- <3> Zero out gradients before each backprop or they'll accumulate
- <4> Forward pass
- <5> Compute loss
- <6> Backpropagation, compute gradients
- <7> Adjust parameters based on gradients
- <8> accumulate batch loss so we can average over epoch

```

[ ]: from torch import optim
from torch import nn

```



```

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(vgg_model.parameters(), # <1>
                      lr=0.001,
                      momentum=0.9)

```

```

[ ]: N_EPOCHS = 10
for epoch in range(N_EPOCHS): # <1>

    epoch_loss = 0.0
    for inputs, labels in trainloader:
        inputs = inputs.to(device) # <2>
        labels = labels.to(device)

        optimizer.zero_grad() # <3>

        outputs = vgg_model(inputs) # <4>
        loss = criterion(outputs, labels) # <5>
        loss.backward() # <6>
        optimizer.step() # <7>

    epoch_loss += loss.item() # <8>
    print("Epoch: {} Loss: {}".format(epoch,
                                       epoch_loss/len(trainloader)))

```

```

Epoch: 0 Loss: 0.8118732511878014
Epoch: 1 Loss: 0.5086276560378075
Epoch: 2 Loss: 0.4192588477998972
Epoch: 3 Loss: 0.36131201485037806
Epoch: 4 Loss: 0.3123959180480242
Epoch: 5 Loss: 0.2794492025240511
Epoch: 6 Loss: 0.25302029708057644
Epoch: 7 Loss: 0.23161081992231308
Epoch: 8 Loss: 0.20988165496785194
Epoch: 9 Loss: 0.19169600319910796

```

```

[ ]: num_correct = 0.0

for x_test_batch, y_test_batch in testloader:

    vgg_model.eval()

    y_test_batch = y_test_batch.to(device)

    x_test_batch = x_test_batch.to(device)

    y_pred_batch = vgg_model(x_test_batch)

    _, predicted = torch.max(y_pred_batch, 1)

```

```

num_correct += (predicted == y_test_batch).float().sum()

accuracy = num_correct/(len(testloader)*testloader.batch_size)

print(len(testloader), testloader.batch_size)

print("Test Accuracy: {}".format(accuracy))

```

625 16

Test Accuracy: 0.899899959564209

```

[ ]: from torch import nn
import torch.nn.functional as F

class LeNet5(nn.Module):
    def __init__(self):
        super(LeNet5, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5) # <1>
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, int(x.nelement() / x.shape[0]))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

device = "cuda" if torch.cuda.is_available() else "cpu"
LeNet_model = LeNet5().to(device=device)

```

```

[ ]: from torch import optim
from torch import nn

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(LeNet_model.parameters(), # <1>
                      lr=0.001,
                      momentum=0.9)

```

```
[ ]: N_EPOCHS = 10
for epoch in range(N_EPOCHS): # <1>

    epoch_loss = 0.0
    for inputs, labels in trainloader:
        inputs = inputs.to(device) # <2>
        labels = labels.to(device)

        optimizer.zero_grad() # <3>

        outputs = LeNet_model(inputs) # <4>
        loss = criterion(outputs, labels) # <5>
        loss.backward() # <6>
        optimizer.step() # <7>

    epoch_loss += loss.item() # <8>
    print("Epoch: {} Loss: {}".format(epoch,
        epoch_loss/len(trainloader)))
```

```
Epoch: 0 Loss: 1.910171902961731
Epoch: 1 Loss: 1.5982797781372071
Epoch: 2 Loss: 1.4726939439201354
Epoch: 3 Loss: 1.3984185062599181
Epoch: 4 Loss: 1.335036766834259
Epoch: 5 Loss: 1.2921974471855164
Epoch: 6 Loss: 1.2550513487911223
Epoch: 7 Loss: 1.2278237766170501
Epoch: 8 Loss: 1.206542158164978
Epoch: 9 Loss: 1.188108439283371
```

```
[ ]: num_correct = 0.0

for x_test_batch, y_test_batch in testloader:

    LeNet_model.eval()

    y_test_batch = y_test_batch.to(device)

    x_test_batch = x_test_batch.to(device)

    y_pred_batch = LeNet_model(x_test_batch)

    _, predicted = torch.max(y_pred_batch, 1)

    num_correct += (predicted == y_test_batch).float().sum()
```

```

accuracy = num_correct/(len(testloader)*testloader.batch_size)

print(len(testloader), testloader.batch_size)

print("Test Accuracy: {}".format(accuracy))

```

625 16

Test Accuracy: 0.616100013256073

```

[ ]: print(LeNet_model)
     print(vgg_model)

```

```

LeNet5(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace=True)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace=True)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace=True)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace=True)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
    (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (13): ReLU(inplace=True)
    (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (15): ReLU(inplace=True)
    (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
    (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (18): ReLU(inplace=True)
    (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (20): ReLU(inplace=True)
    (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

        (22): ReLU(inplace=True)
        (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (25): ReLU(inplace=True)
        (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (27): ReLU(inplace=True)
        (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (29): ReLU(inplace=True)
        (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))
    (classifier): Sequential(
      (0): Linear(in_features=25088, out_features=4096, bias=True)
      (1): ReLU(inplace=True)
      (2): Dropout(p=0.5, inplace=False)
      (3): Linear(in_features=4096, out_features=4096, bias=True)
      (4): ReLU(inplace=True)
      (5): Dropout(p=0.5, inplace=False)
      (6): Linear(in_features=4096, out_features=10, bias=True)
    )
  )
)

```